



# Databricks Assistant



Boost productivity with your  
**context-aware AI** assistant



---

Caio Moreno | Solutions Architect  
Marine Tranquard | Account Executive



**92%**

U.S.-based  
developers are  
already using **AI  
coding tools** both  
in and outside of  
work.

Source: [Beet Roots](#)

**70%**

say AI coding tools will  
offer them an  
**advantage at work**  
and cite better code  
quality, completion  
time, and resolving  
incidents

Source: [Github blog 2024](#)

Enterprises want to  
**empower their data practitioners**  
with natural language tools

# AI assistants don't work on enterprise data

**Unable to locate the right data assets**



AI doesn't know where to find the relevant tables, pipelines, queries, and dashboards

**Only generate generic responses**



AI doesn't understand your unique data structures, usage patterns, or business terminology

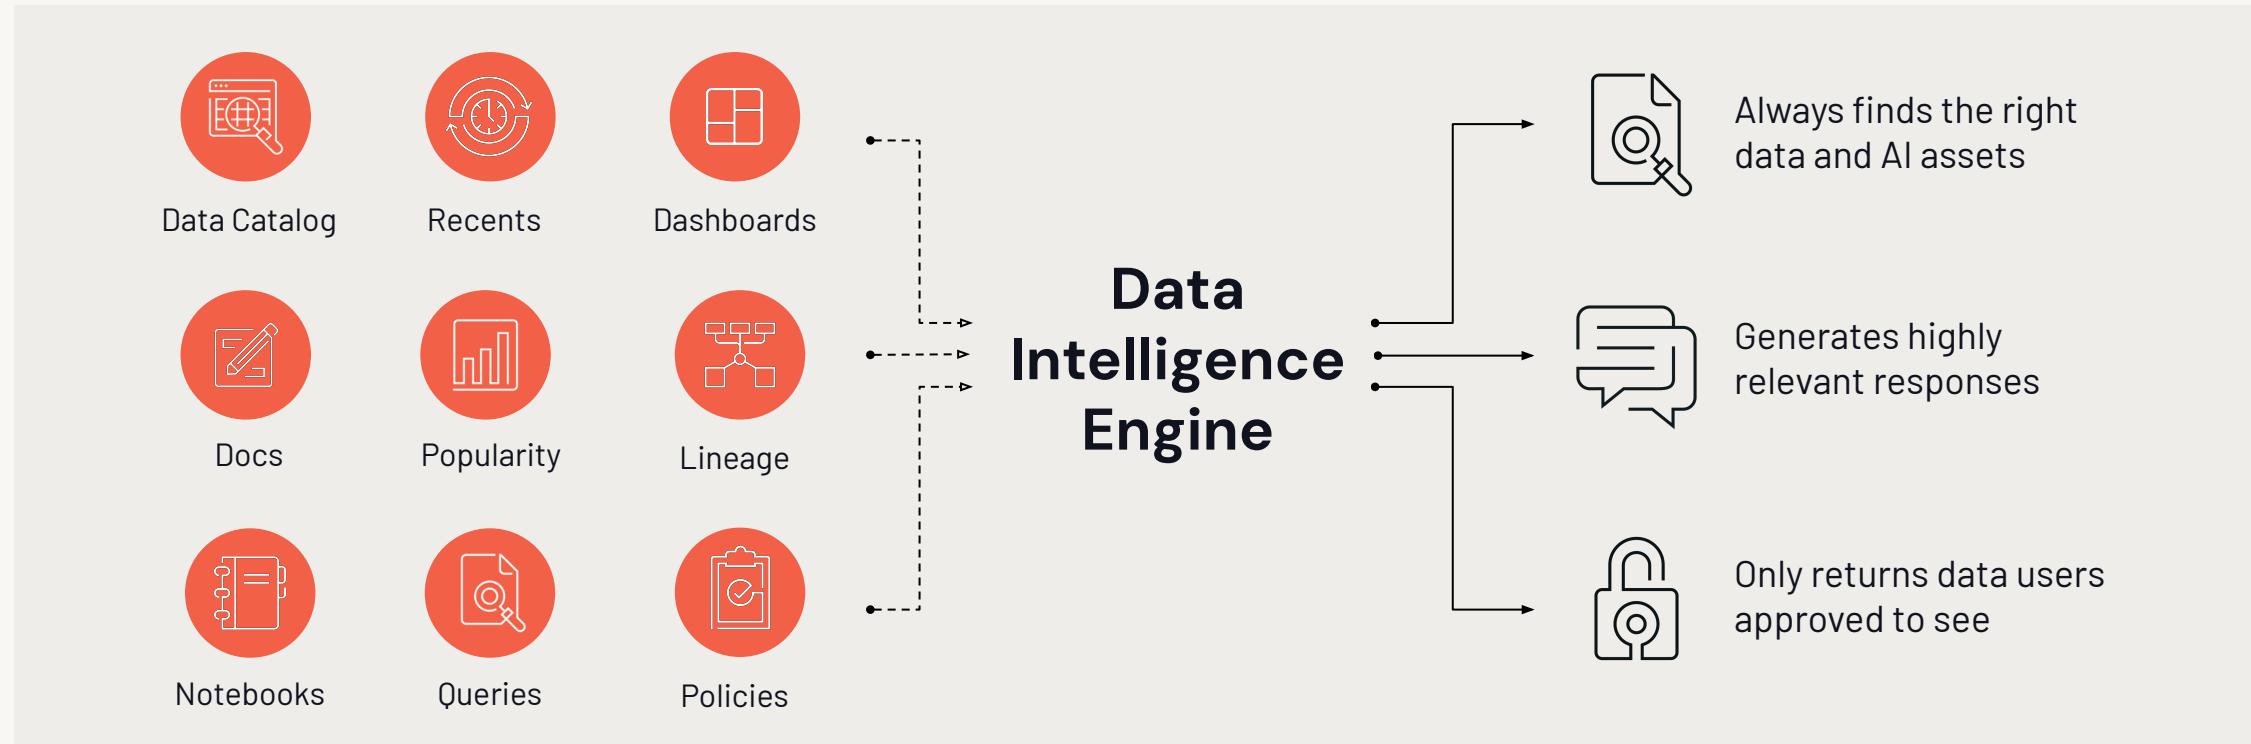
**Don't comply with access policies**



AI doesn't know which users and groups are allowed to see which schemas, tables, and columns

## Generalized LLMs lack data intelligence

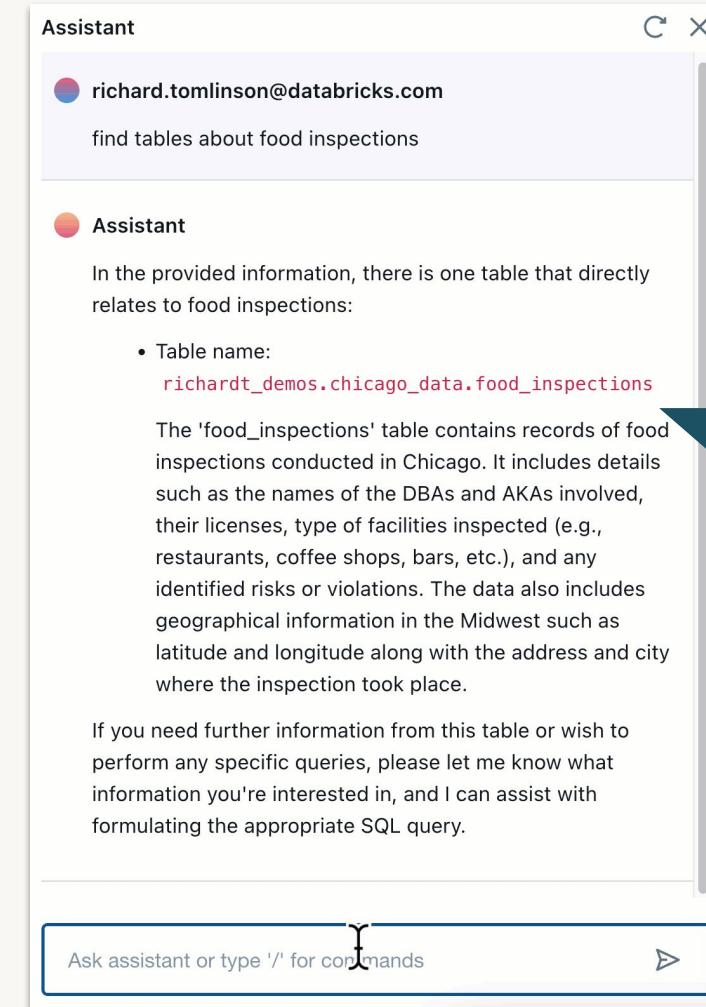
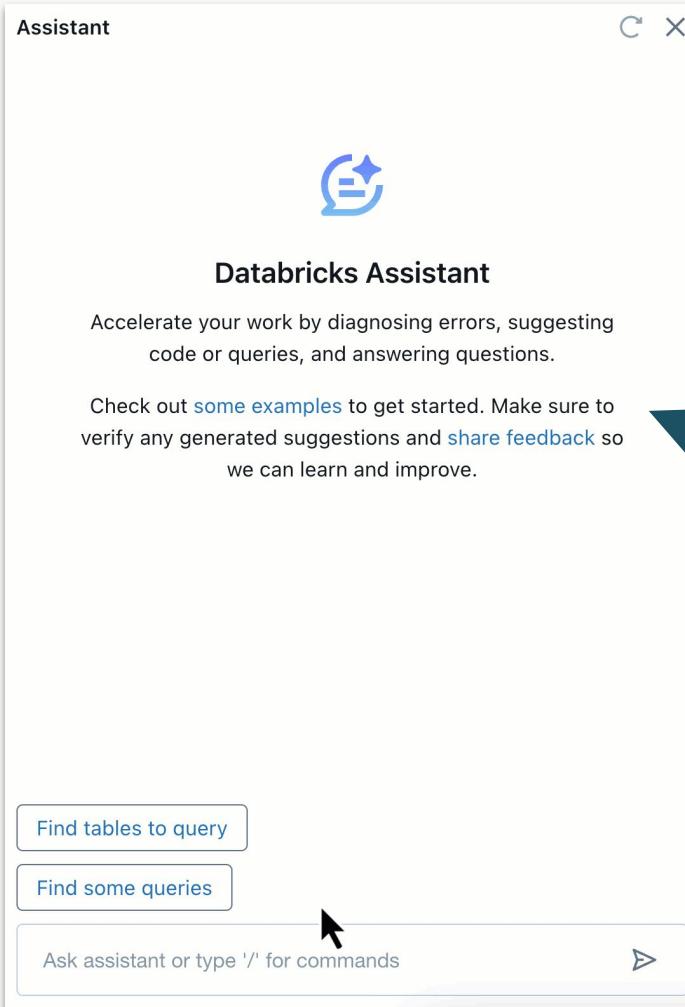
# Data intelligence gives AI the unique context of your business



Data intelligence engine uses AI to learn the signals in your data platform

AI assistant gains awareness of your <sup>5</sup> unique data and business context

# Example of a context-aware AI assistant



# Databricks Assistant

Boost productivity with your context-aware AI assistant



- Generate, explain, and fix code with natural language
- AI assistance in every user experience
- Powered by DatabricksIQ for highly relevant answers

A screenshot of the Databricks Notebook interface. On the left is the navigation sidebar with options like New, Workspace, Recents, Catalog, Workflows, Compute, SQL, SQL Editor, Queries, Dashboards, Alerts, Query History, SQL Warehouses, Data Engineering, Job Runs, Data Ingestion, Delta Live Tables, Machine Learning, Playground, Experiments, and Features. The main area shows an "Untitled Notebook 2024-01-26 12:27:52" in Python. A central panel titled "Assistant" contains the "Databricks Assistant" logo and text: "Accelerate your work by diagnosing errors, suggesting code or queries, and answering questions. Check out some examples to get started. Make sure to verify any generated suggestions and share feedback so we can learn and improve." Below this are buttons for "Apply machine learning", "Run data summarization", and "Find tables to query". A tooltip box is overlaid on the interface, containing the text: "I have a CSV file in the UC volume /Volumes/richardt\_demos/chicago\_data/restaurant\_data and I need some pyspark code to write it to a UC managed table called richardt\_demos.chicago\_data.restaurant\_inspections. please infer the schema on the CSV file." A cursor arrow points towards this tooltip. The top right of the notebook window shows user information: "E2 Dogfood", "richard.tomlinson@databricks.com", and various sharing and scheduling buttons.



Over 100K users assisted every month!



# Generate, explain, and fix code with natural language



- Takes requests in natural language and creates code snippets
- Applies details from code cells, libraries, runtime, and more to improve accuracy
- Explains, diagnoses, and fixes issues from within a cell

A screenshot of the Databricks SQL Editor interface. On the left is a dark sidebar with various navigation options like Workspace, Recents, Data, Workflows, Compute, and several sections under SQL (SQL Editor, Queries, Dashboards, Alerts, Query History, SQL Warehouses), Data Engineering (Delta Live Tables), Machine Learning (Experiments, Features, Models, Serving), and General (Clusters, Storage, Metrics, Experiments, General). The main area is titled 'Assistant'. It contains a message: 'The Assistant can accelerate your work by diagnosing errors, suggesting code or queries, and answering questions. Mistakes are possible, so be sure to check facts and share your feedback.' Below this is a text input field with placeholder text 'Start with picking the tables to query?' and a button 'Enter your message'. At the bottom, there's a scrollable list of AI-generated suggestions. A large blue cube icon with a minus sign is positioned on the right side of the main window. The status bar at the bottom shows 'Run a query to show the results' and some other status information.

# AI assistance in every user experience



The screenshot displays four different Databricks products with callout bubbles highlighting the AI Assistant feature:

- Databricks Assistant in SQL editor:** Shows the AI Assistant in the bottom-left corner of the SQL editor interface.
- Databricks Assistant in Notebooks:** Shows the AI Assistant in the top-left corner of the Notebook interface.
- Databricks Assistant in Lakeview dashboards:** Shows the AI Assistant in the bottom-right corner of the Lakeview dashboard interface.
- Databricks Assistant in Help:** Shows the AI Assistant in the top-right corner of the Help interface.

**Boosts productivity for all users**

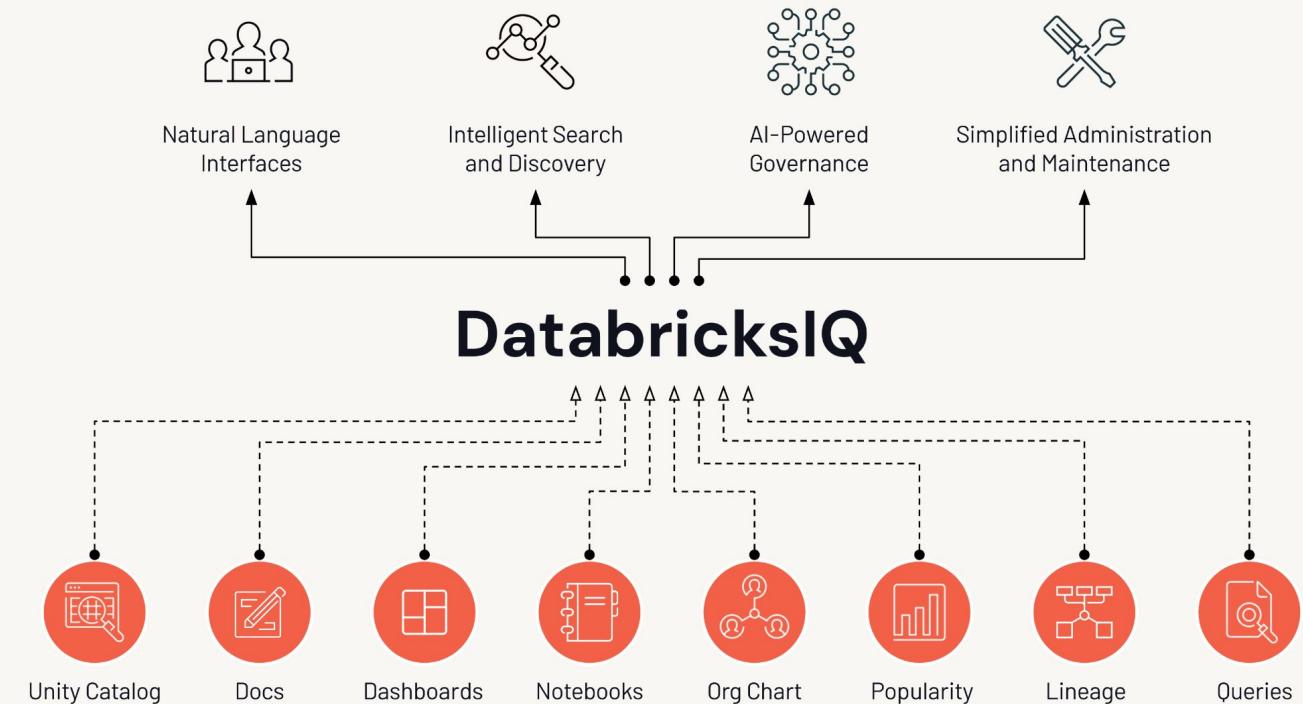
Data Scientists | Data Engineers | SQL Analysts | Business Users

# Powered by DatabricksIQ

The data intelligence engine that understands your business



- Apply AI to your data to provide relevant, accurate, and secure answers
- Enables natural language interfaces for everyone
- Uses predictive models to automatically optimize your data platform



*“How many DBUs were there in Europe last quarter?”*



# Example: Assistant with and without DatabricksIQ

**Without**  
DatabricksIQ

The AI is new, so surprises and mistakes are possible.  
Make sure to check the facts, and share feedback.

**user@databricks.com**  
What is our total revenue in Europe?

**Databricks**  
Here is a SQL query for **revenue** in **Europe**:

```
sql
SELECT
    sum(sales_amount)
FROM
    dbdemos.sales.orders_gold
WHERE
    salesTerritory = 'Europe'
```

Ask a question... ▶

**With**  
DatabricksIQ

Powered by  
**Databricks IQ**

The AI is new, so surprises and mistakes are possible.  
Make sure to check the facts, and share feedback.

**user@databricks.com**  
What is our total revenue in Europe?

**Databricks**  
Here is a SQL query for **revenue** in **European sales territories**:

```
sql
SELECT
    sum(sales_amount)
FROM
    dbdemos.sales.orders_gold
WHERE
    -- European sales territories
    salesTerritory IN (
        'EMEA Northern',
        'EMEA Southern'
    )
    -- Exclude internal orders
    AND customerId != 'DatabricksInternal'
```

Ask a question... ▶



# GPT Misses

Model: GPT-4

MO

Where can I find data on the number of DBU's generated by Nephos workloads?



I'm not aware of a specific source that tracks the number of "DBUs" (Database Units) generated by "Nephos" workloads, as Nephos isn't a known technology or company related to databases, at least as of my knowledge cutoff in September 2021. It's possible that you're referring to a new technology, service, or company that emerged after that date, or it could be a typo.

If you're looking for data on a specific database technology or service, please provide more context or clarify the terms you're using, and I'll do my best to help you find the information you're looking for.



Model: GPT-4

MO

Give me Photon speed metrics.



The speed of a photon is the speed at which light travels through a vacuum. This is commonly referred to as the speed of light, and is a fundamental constant of nature. The speed of light is approximately:

1. 299,792,458 meters per second (m/s)
2. 186,282 miles per second (mi/s)
3. 670,616,629 miles per hour (mi/h)
4. 1,080,000,000 kilometers per hour (km/h)

It's important to note that the speed of light can vary when it travels through different media, such as air, water, or glass. In these cases, the speed of light will be lower than its speed in a vacuum. However, when discussing the speed of photons or the speed of light, it is generally assumed to be in a vacuum unless otherwise specified.



# Context is king

“How many DBUs were there in Europe in the last quarter?”

- **Well-documented Code and Queries:** Are there samples we can use as examples?
- **Popularity:** What are the most used tables across my company?
- **Favorites:** Which tables I have used recently? Frequently?
- **Documentation:** Are there relevant wikis and docs for interacting with this data?
- **Organization:** Who accessed these tables? Do I frequently work with them?
- **Lineage:** When was this table last updated?
- **Dashboards:** what are the highly-vetted queries powering frequently-used charts?
- **UC Descriptions and Tags:** What are the tables certified by my data team?



## Personalized Gen-AI Experiences

### Databricks Assistant

Code gen, text-to-sql, fix my code

### Data Asset Discovery

Find me the right table

### AI Documentation

Enrich metadata

### Gen-AI Platform

(LLMs, Vector Index, RAG)

### Semantic Knowledge Graph

(Data, People, Activity)

### Unity catalog

One security and governance model for all data and AI across the organization

### Delta Lake

One platform to store and manage all structured, semi-structured, and unstructured data



### Cloud Data Lake

All Raw Data  
(Logs, Texts, Audio, Video, Images)



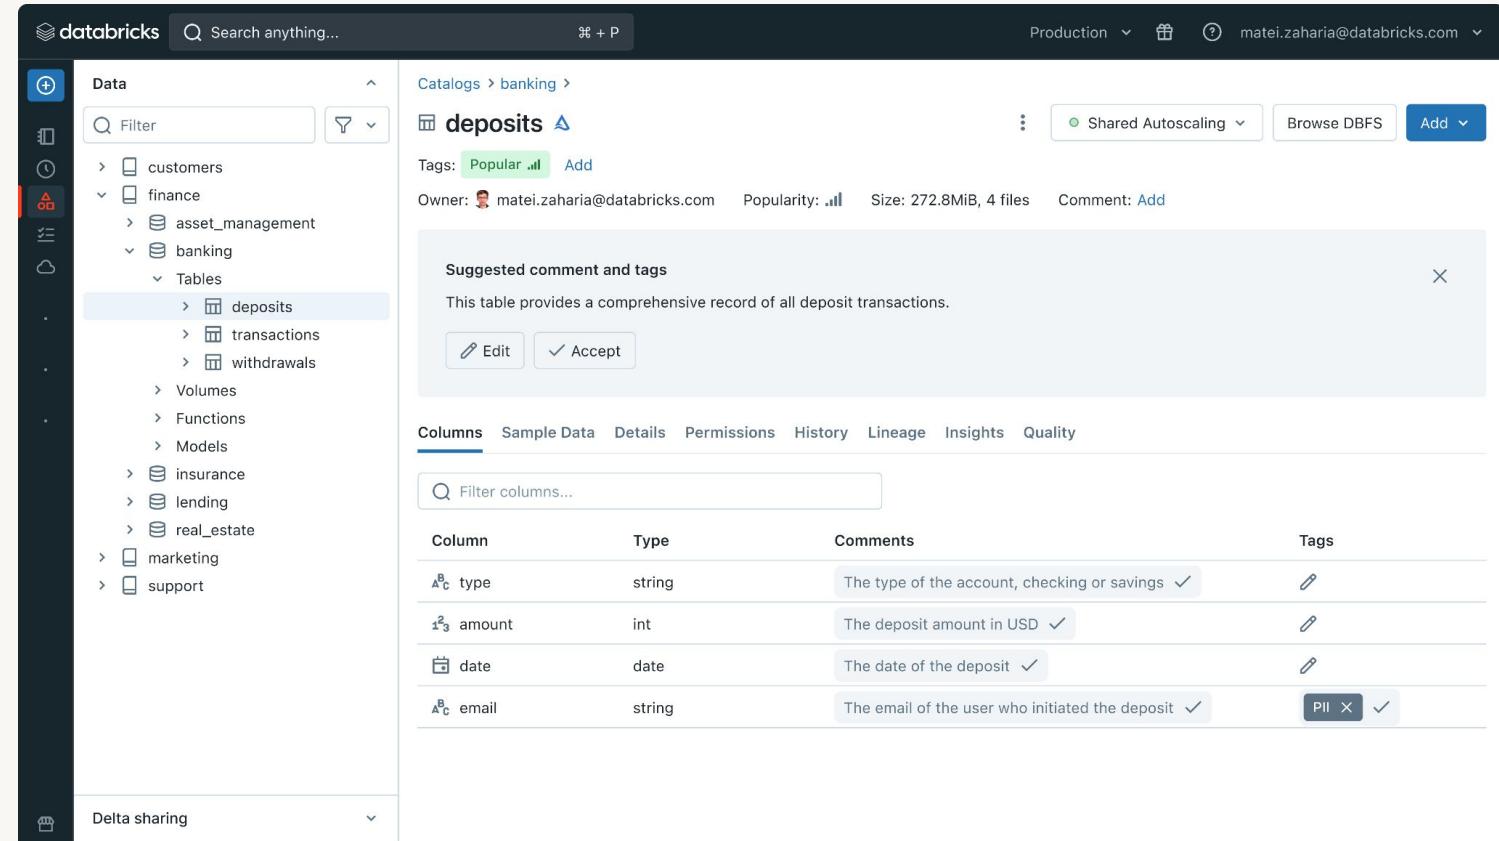
**DatabricksIQ**

# Personalized Gen-AI powered by DatabricksIQ



# Unity Catalog Data Documentation

- Auto-generate concise and informative **table and column comments** for Unity Catalog
- Document your backlog of data assets with **missing documentation** in minutes



The screenshot shows the Databricks Unity Catalog interface. On the left, there's a sidebar with a search bar and a navigation tree under 'Data'. The tree shows categories like 'customers', 'finance' (which is expanded), 'asset\_management', 'banking' (also expanded), and 'marketing', 'support'. Under 'banking', 'Tables' is expanded, showing 'deposits', 'transactions', 'withdrawals', 'Volumes', 'Functions', 'Models', 'insurance', 'lending', 'real\_estate', 'marketing', and 'support'. The 'deposits' table is selected. The main panel shows the 'deposits' table details. It includes a 'Tags' section with 'Popular' and an 'Add' button. Below that is a 'Suggested comment and tags' section with a message: 'This table provides a comprehensive record of all deposit transactions.' There are 'Edit' and 'Accept' buttons. At the bottom, there's a table with columns 'Column', 'Type', 'Comments', and 'Tags'. The columns listed are 'type' (string, checked PII), 'amount' (int, checked PII), 'date' (date, checked PII), and 'email' (string, checked PII).

Column	Type	Comments	Tags
type	string	The type of the account, checking or savings ✓	✓
amount	int	The deposit amount in USD ✓	✓
date	date	The date of the deposit ✓	✓
email	string	The email of the user who initiated the deposit ✓	PII ✓



# Databricks Assistant

- **AI-powered Authoring Assistant** integrated into the notebook, file and SQL editors.
- Generate, Fix, and Explain code and queries

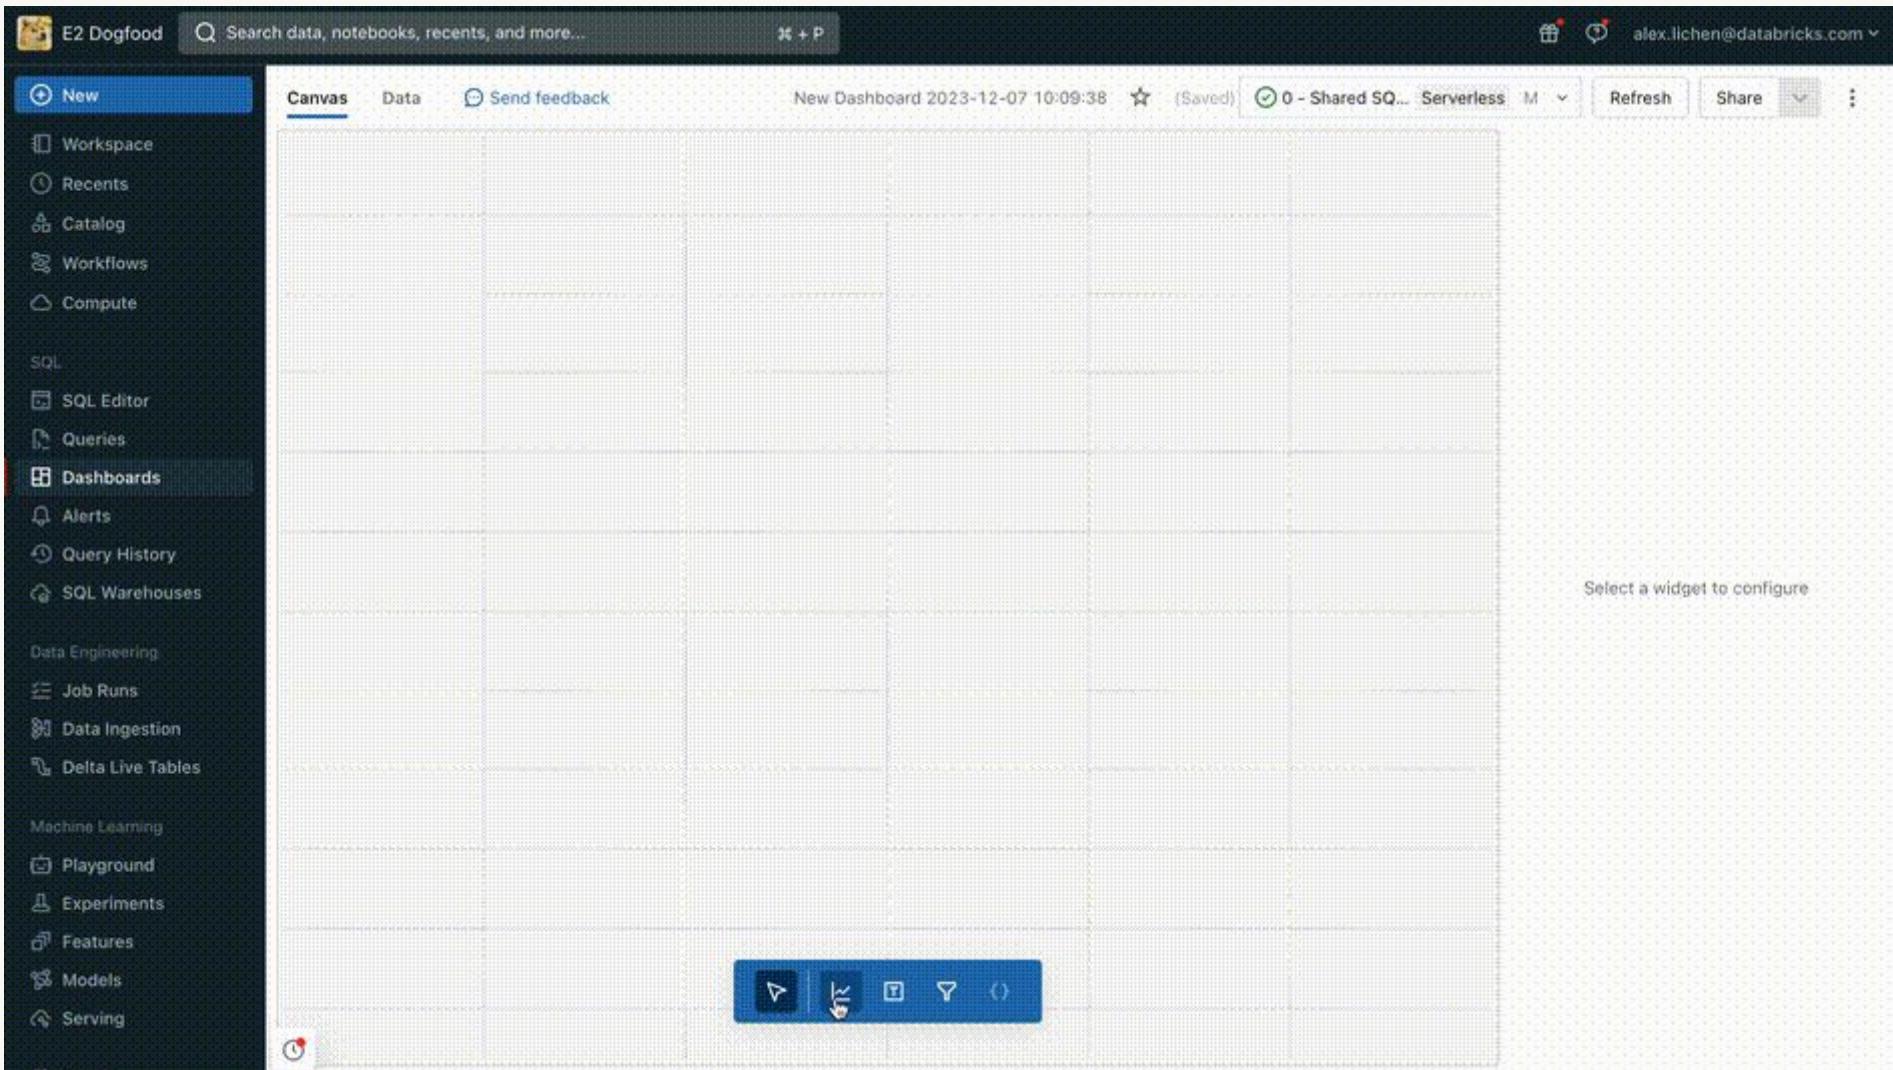
The screenshot shows the Databricks Assistant interface within a notebook titled "Databricks Assistant Demo" in Python. The sidebar on the left contains links for Workspace, Recents, Data, Workflows, Compute, SQL (SQL Editor, Queries, Dashboards, Alerts, Query History, SQL Warehouses), Data Engineering (Delta Live Tables), Machine Learning (Experiments, Features, Models), and Serving. The main area is titled "Assistant" and contains a brief introduction: "The Assistant can accelerate your work by diagnosing errors, suggesting code or queries, and answering questions. Mistakes are possible, so be sure to check facts and share your feedback. (More info)". Below this is a message input field with the placeholder "Start with picking the tables to query?" and a "Enter your message" button. To the right is a results table with two rows:

ID	Created At	Last Run At
119	2016-02-27 22:26:56.000	2016-02-27 22:31:25.000
120	2016-01-31 17:52:43.000	2016-01-31 17:57:44.000

Below the table, a note states: "Command took 1.83 seconds -- by jim.allenwallace@databricks.com at 7/30/2023, 3:12:55 PM on 13.x Shared". A tooltip provides keyboard shortcuts: Shift+Enter to run, Shift+Ctrl+Enter to run selected text, and Option+Shift+Space to suggest code.



# Assistant for Lakeview



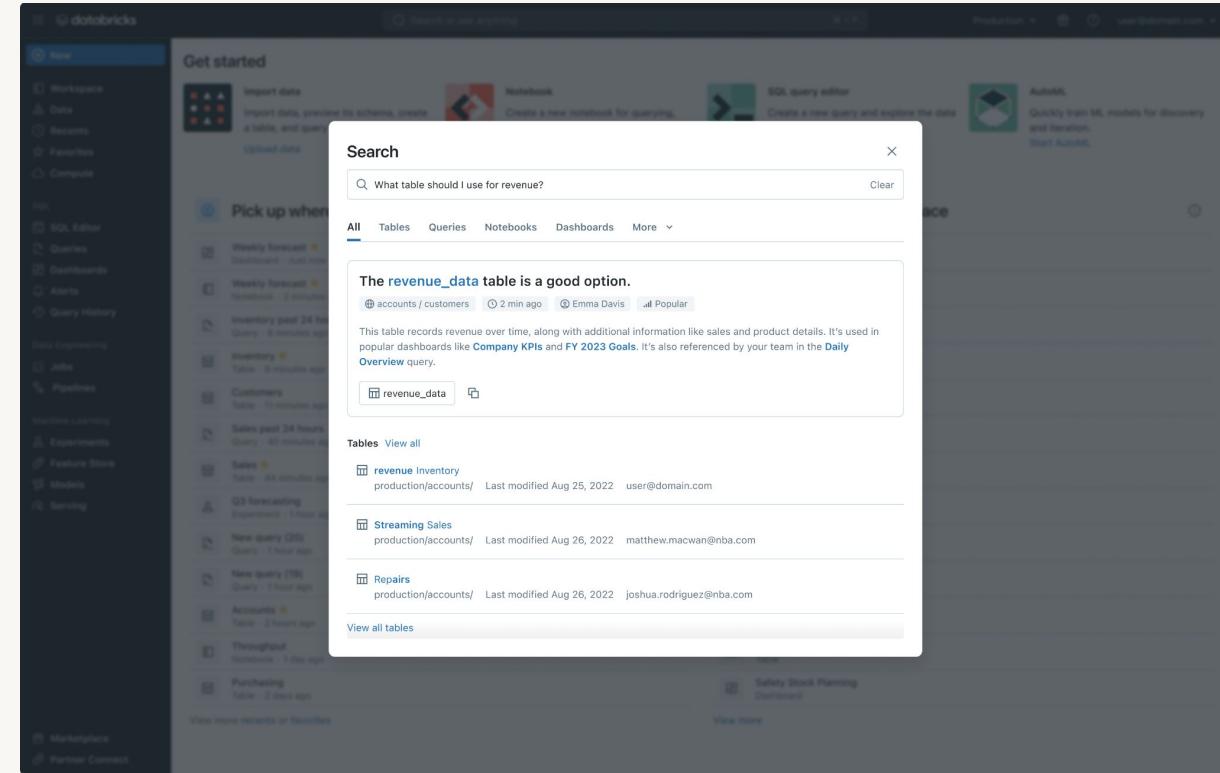
# Intelligent Search

## Contextual and accurate search and knowledge card

Accurate semantic search using enterprise knowledge graph

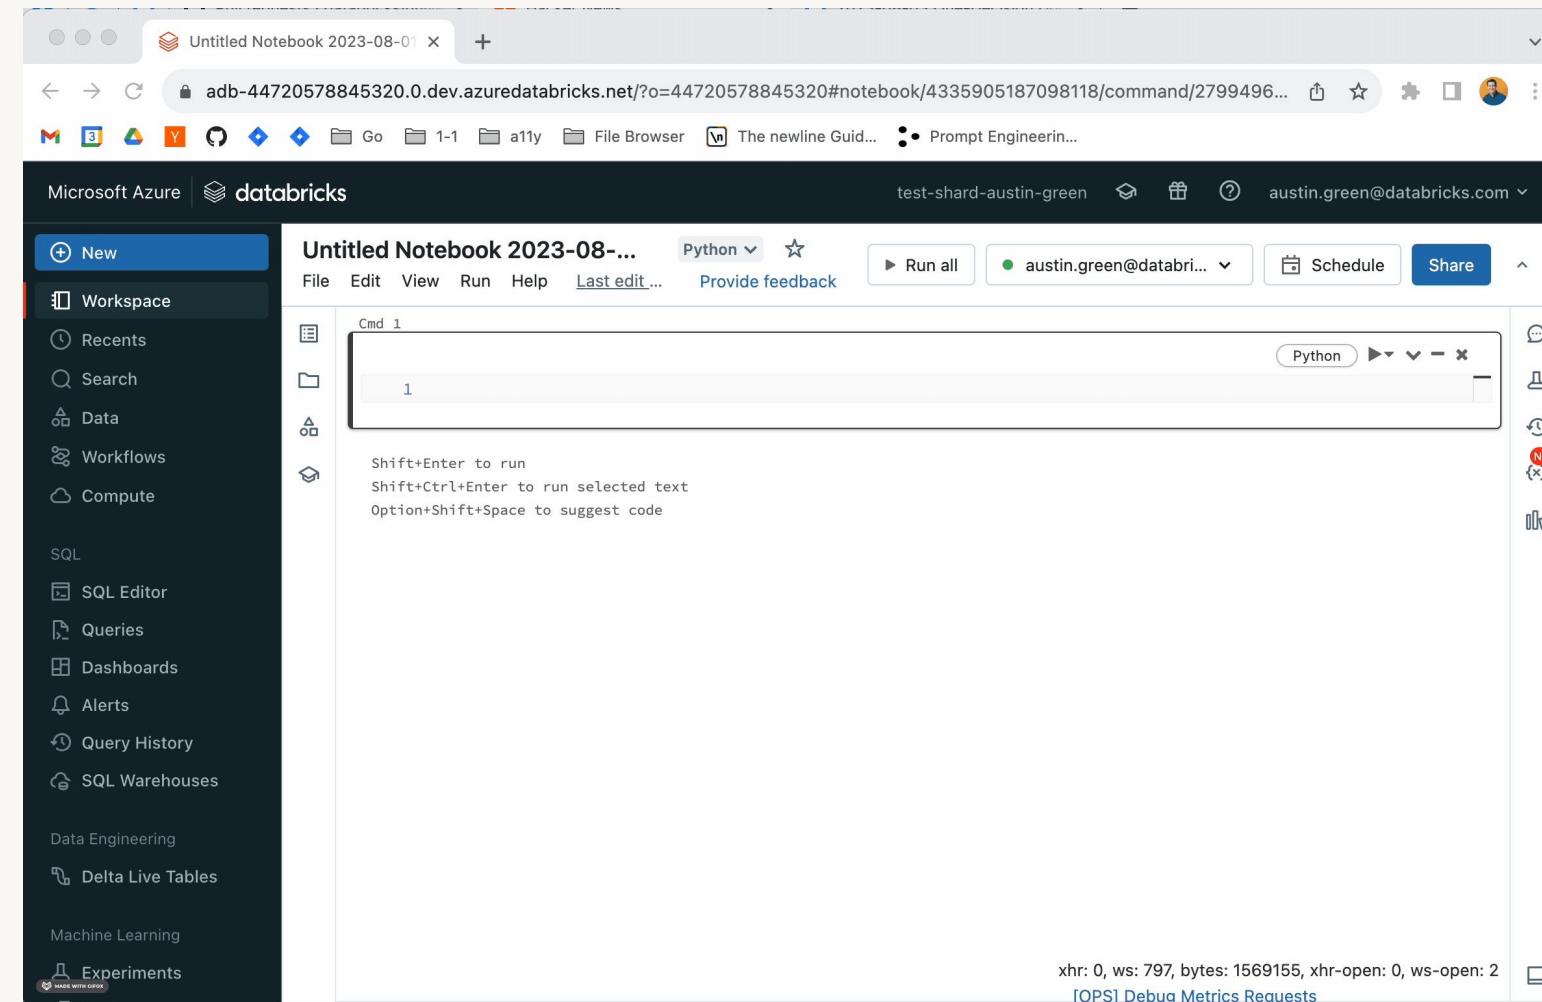
Summary knowledge card with details to help quickly get to what you are looking for

LLM-powered Natural language understanding



# AI Code Suggestions-As-You-Type

- “Ghost text” single and multi-line suggestions that automatically appear as you type.
- Completions available for Python, Scala, R and SQL.



# Project Genie

**Enable business users to interact with data with LLM-powered Q&A**

Ask questions in natural language and receive answers in text and visualizations

Curate dataset-specific experiences with custom instructions

Powered by Databricks SQL & DatabricksIQ



A screenshot of the Project Genie interface. On the left, there's a card for "EMEA Sales Analysis Data" with options to "Edit" or "Share". Below the card are links for "Chats", "Tables", and "Instructions". At the bottom of this section is a "Back to list" button. To the right of this card is a large white area with a "New chat" button at the top. Below it, a message says "Conversation history is currently disabled." Underneath are "Quick actions" buttons for "Example questions", "Explain data set", and "Surprise me". At the bottom of the interface is a text input field with the placeholder "Name all the unique stores with their footfall" and a red circular icon with "VA" and the name "Vidhi Agrawal" next to it.

Now in Private Preview



# Centralized AI Governance

Unity Catalog provides a secure, governed collaboration layer

## Available Now

- Enable/Disable in account or per-workspace
- Assistant only uses table and comment description; it doesn't look at row-level data
- Uses user permissions (e.g., it does not send metadata relating to tables that the user does not have permission to see)

## H1 2024

- Assistant integrates with audit logs so you can see usage in your workspace
- [Potential] Prioritize tables for the Assistant to use.
- [Potential] Human-in-the-loop curation/validation



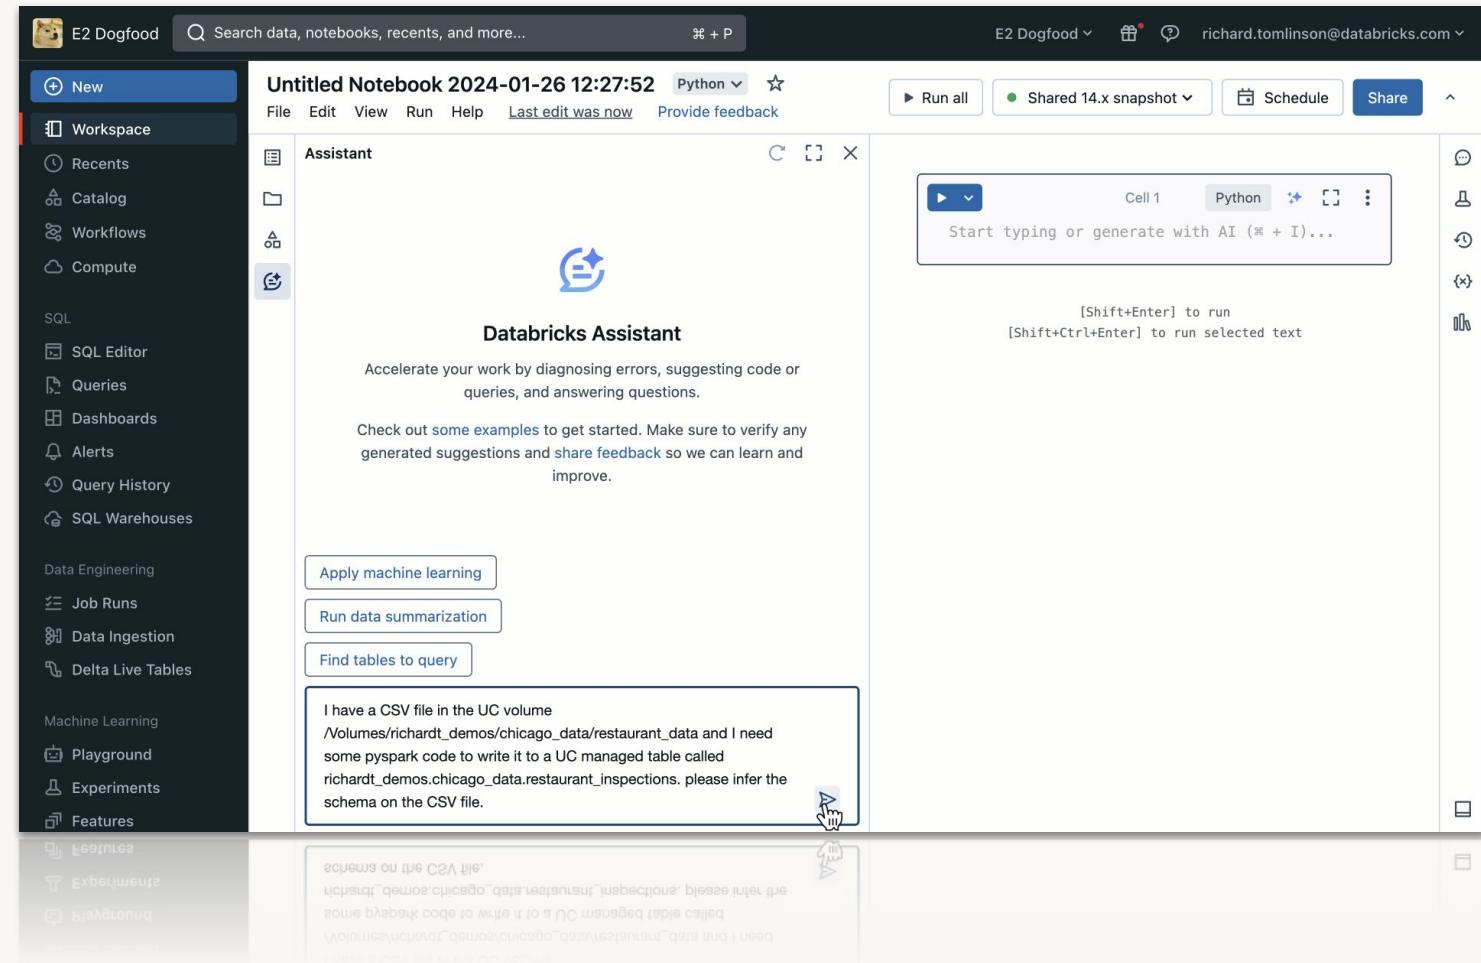
# Databricks Assistant

Your context-aware AI assistant.

Automatically generates SQL and Python, explains complex code, and fixes issues.

AI assistance in every user experience; Notebooks, SQL & File editors, Lakeview, Help, and more.

Powered by DatabricksIQ to ensure highly relevant responses based on your data and usage.



# Example: Search with and without DatabricksIQ

The image displays two side-by-side search interfaces. The left interface, labeled "Without DatabricksIQ", shows a basic search bar and navigation links for All, Tables, Notebooks, Jobs, Queries, and More. It lists a notebook named "02-forecast-billing-tables" with a SQL code snippet. The right interface, labeled "With DatabricksIQ", shows a search bar and navigation links. It lists a table named "Nephos KPIs" with a "Popular .sql" link. Both interfaces include a "Powered by DatabricksIQ" logo.

**Without DatabricksIQ**

Search

Where can I see serverless usage?  Clear

All Tables Notebooks Jobs Queries More

Notebooks [View all notebooks](#)

02-forecast-billing-tables

/Users/mason.clark@databricks.com/SystemTables/billing-tables ian.sharpe@databricks.com modified: 6/22/2023

SQL

```
1 %sql
2 select * from detailed_billing_forecast where sku = 'ALL' limit 100
```

**With DatabricksIQ**

Search

Where can I see serverless usage?  Clear

Powered by **DatabricksIQ**

All Tables Notebooks Jobs Queries More

Tables [View all tables](#)

Nephos KPIs [Popular .sql](#)

Key performance indicators (KPIs) in real-time powered by **Nephos**. Table includes **DBUs**, and other key...

eve.adebayo@.../Demo/KPI eve.adebayo@databricks.com modified: June 9, 2023 Frequent users: S D E +2

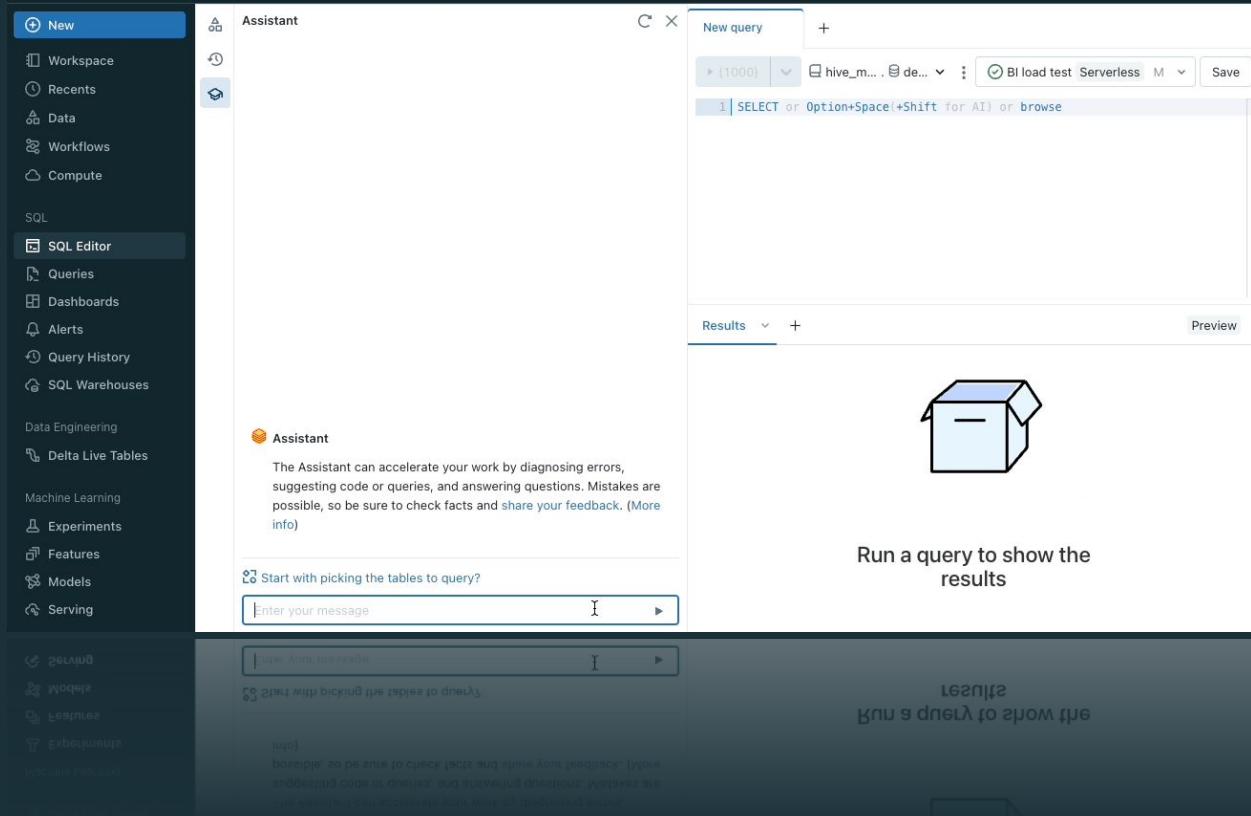
# What can you do with Databricks Assistant?

# Generate SQL or Python Code

Enter natural language requests to generate code

Use natively from Notebooks, SQL editor, and file editor

Copy code or add it directly into a cell



# Autocomplete code or queries



- Generate code from comments inside a notebook cell or SQL editor

The screenshot shows the Databricks Assistant Demo - Code Completion interface. The left sidebar includes options like New, Workspace, Recents, Data, Workflows, Compute, SQL, SQL Editor, Queries, Dashboards, Alerts, Query History, and SQL Warehouses. The main area is titled "Databricks Assistant Demo - Code Completion" and shows a Python notebook cell labeled "Cmd 1". The cell contains the following code:1 # Write code to reverse a string. Then display then the reversed string.

Below the cell, there are three lines of instructions: "Shift+Enter to run", "Shift+Ctrl+Enter to run selected text", and "Option+Shift+Space to suggest code". The top navigation bar includes File, Edit, View, Run, Help, and a "Provide feedback" button. The top right features status indicators for two jobs (green circles), an "Interrupt" button, a "13.x Shared" dropdown, a "Schedule" button, and a "Share" button. The bottom navigation bar includes "secondJobID" and "thirdJobID".

# Transform code



- Always use the best language for the current task
- Utilize alternative frameworks seamlessly

A screenshot of the Databricks Assistant Demo interface. The left sidebar shows various workspace sections like Workspace, Data, Workflows, Compute, SQL, and Data Engineering. The main area is titled "Databricks Assistant Demo" and shows a table of logs with columns for ID, Date, and Time. A tooltip indicates a command took 1.83 seconds. Below the table, there's an "Assistant" section with a message about its capabilities and a text input field for interacting with it.

ID	Date	Time
119	2016-02-27	22:26:56.000
120	2016-01-31	17:52:43.000

Command took 1.83 seconds -- by jim.allenwallace@databricks.com at 7/30/2023, 3:12:55 PM on 13.x Shared

Shift+Enter to run  
Shift+Ctrl+Enter to run selected text  
Option+Shift+Space to suggest code

**Assistant**  
The Assistant can accelerate your work by diagnosing errors, suggesting code or queries, and answering questions. Mistakes are possible, so be sure to check facts and share your feedback. (More info)

Start with picking the tables to query?

Enter your message

# Explain code or query



- Get a description of code in clear, concise language
- See what data and tables are being used without needing to leave the editor

The screenshot shows the Databricks SQL Editor interface. On the left is the sidebar with various notebooks and workspace items. The main area has a title bar "Assistant" and a message from "jim.allenwallace@databricks.com": "write a query to return the 10 most expensive fares from main.nyctaxi.trips". Below this, another "Assistant" box contains a sample SQL query:

```
sql
SELECT fare_amount
FROM main.nyctaxi.trips
ORDER BY fare_amount DESC
LIMIT 10;
```

A detailed explanation follows: "This query selects the `fare_amount` column from the `main.nyctaxi.trips` table, orders the results in descending order based on `fare_amount`, and returns only the top 10 rows." To the right, the results of the query are displayed in a table:

#	fare_amount
1	275.00
2	260.00
3	188.00
4	130.00
5	115.00
6	105.00
7	105.00
8	105.00
9	105.00
10	105.00

At the bottom, there's a preview section showing a smaller table with the same data.

# Fix issues fast



- Identify errors and see an explanations
- Receive recommended fixes
- Update code with suggested snippets to resolve the issue

The screenshot shows the Databricks Assistant Demo interface. On the left is a dark sidebar with various navigation options like Workspace, Data, Workflows, Compute, SQL, and Machine Learning. The main area has a title bar "Databricks Assistant Demo" and a tab "Python". Below the title bar are buttons for "Run all", "13.x Shared", "Schedule", and "Share". A "Cmd 1" tab is open, showing Python code:

```
1 from pyspark.sql.functions import col
2
3 # create a dataframe with two columns: a and b
4 df = spark.range(5).select(col('id').alias('a'), col('id').alias('b'))
5
6 # try to select a non-existing column c
7 df.select(col('c')).show()
```

A circled button in the top right corner of the code editor has a hand icon over it, likely indicating a suggestion or action. Below the code editor, there's a section titled "Assistant" with a description: "The Assistant can accelerate your work by diagnosing errors, suggesting code or queries, and answering questions. Mistakes are possible, so be sure to check facts and share your feedback. (More info)". There's also a message input field: "Start with picking the tables to query?" and "Enter your message". At the bottom, there's some small text about Databricks features.



# Use natural language to visualize data



- Ask questions in the language of your business
- Add visuals to dashboards and share across the organization
- Empower non-technical users to self-serve

A screenshot of the Databricks Canvas interface. The top navigation bar includes the Databricks logo, a search bar, and various dashboard management buttons like 'Draft', 'Publish', and 'Share'. The main area is titled 'New Dashboard 2024-03-06 10:14:06' and shows a 5x5 grid for placing widgets. A tooltip 'Select a widget to configure' is visible on the right. At the bottom of the grid is a toolbar with icons for adding a new row or column, and the bottom right corner shows a small preview of another dashboard.

A screenshot of the Databricks Data interface, showing a similar 5x5 grid structure to the Canvas interface. It appears to be a workspace for managing data assets or notebooks. A toolbar with icons for adding a new row or column is visible at the bottom right.



# Get help and support

- A conversational interface for product documentation and support
- Get accurate answers for Databricks technologies like Spark, Delta Live Tables, DBSQL, and more.
- Improved flow for creating support tickets

The screenshot shows the Databricks interface with a focus on the 'Help' feature. On the left, there's a sidebar with a blue circular icon containing a white question mark, followed by the word 'Help'. Below this, there are two main sections: 'Popular' and 'Recent'. The 'Popular' section lists several notebooks and dashboards, each with a small icon and a link. The 'Recent' section shows a list of tables from a database. To the right of the sidebar, the main dashboard displays various Databricks services: 'SQL query editor' (with a 'Create query' button), 'AutoML' (with a 'Start AutoML' button), and a 'Help' panel. The 'Help' panel has a 'Welcome to Help' message, a user profile for 'ben.stallworth@databricks.com', a question about 'Unity Catalog', and a detailed answer from 'Databricks' explaining Unity Catalog's features. It also includes a question from 'ben.stallworth@databricks.com' about getting help and an answer from 'Databricks' directing them to Azure Support or the Databricks Help Center. At the bottom, there's a search bar and a 'Log out' button.



Customers ❤️ Databricks Assistant

“

**A groundbreaking AI-powered tool that transforms coding productivity**

Saptagiri Kintali, Morgan Stanley



“

**You can't take this away from me!**

Mike Lavina, 84.51



“

*“For someone that's competent, it's extremely good at accelerating development”*

—Luke Woolley, SSE Business Energy



“

**I was able to code 200+ lines of robust code in a language I've never coded before**

Josue A. Bogran, Kythera Labs



“

**This AI-based companion is set to reshape the way we code and interact with our lakehouse.**

Jeroen Roosen, Intellus Group



“

**The introduction of the Databricks Assistant has made it easier for our user base to improve their skill set.**

Nicholas Heier, General Motors



*“This cutting-edge AI companion has revolutionized my data analysis journey, simplifying complex tasks and accelerating productivity”*

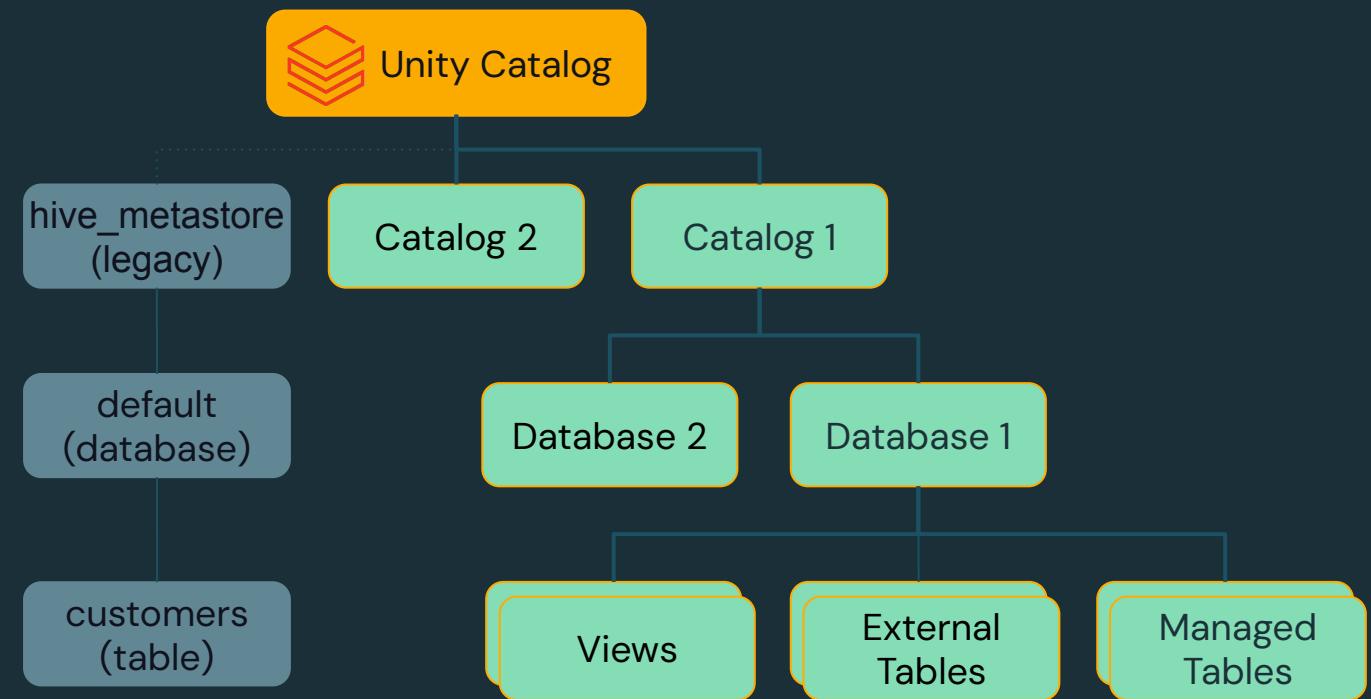
—Byron Exaporriton, ABN AMRO

# Leverages Unity Catalog to provide personalized responses

Applies table schema to locate the right data

Knows popular tables and joins to pick relevant tables

Uses descriptions and tags to understand content



# Available natively in Databricks Notebooks, SQL editor, and file editor

Increases productivity for every user persona :

- CDOs / CTOs
- Data Analysts
- Data Scientists
- Data Engineers





# Databricks Assistant

## Architecture & Security

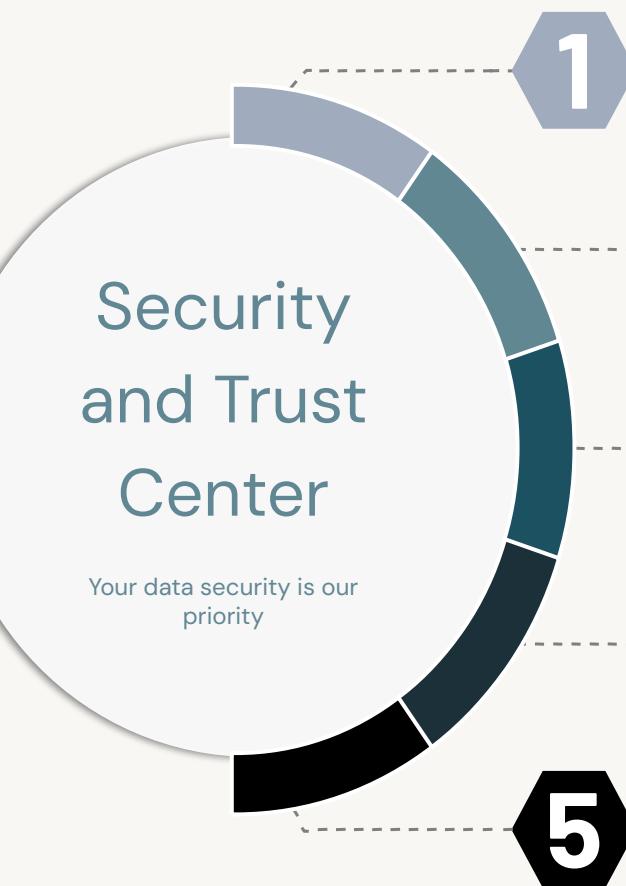
---

Last updated April 2024



# Databricks Security Summary

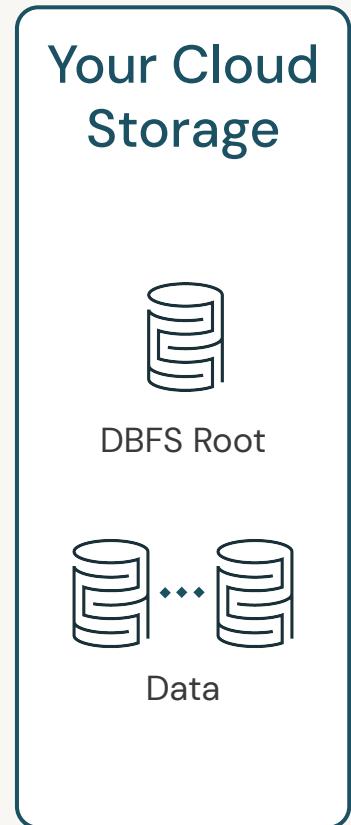
Databricks Security & Trust Center provides an overview of:



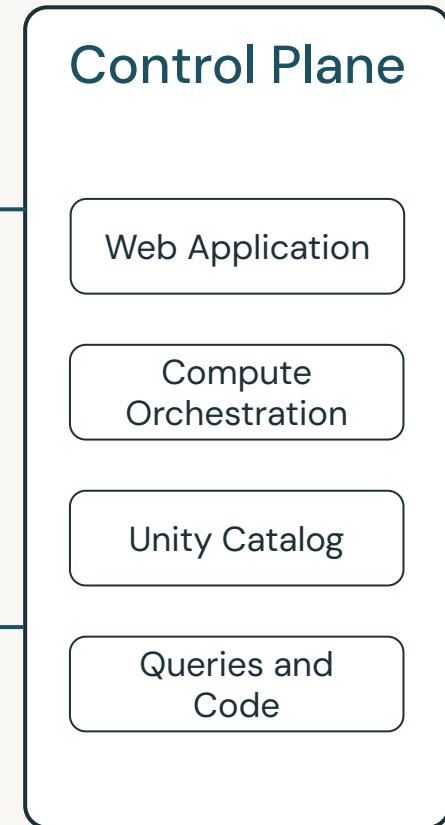
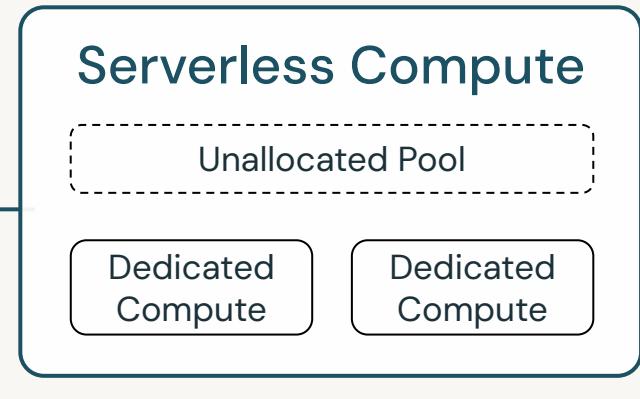
The screenshot shows the homepage of the Databricks Security and Trust Center. The header includes the Databricks logo and navigation links for Platform, Solutions, Learn, Customers, Partners, Company, Try Databricks, Watch Demos, Contact Us, and Login. The main content area has a dark background with white text. It features a large heading "Security and Trust Center" and a sub-heading "Your data security is our priority". Below this are sections for "Trusted Platform", "Security Features", "Compliance", and "Privacy", each with an icon and a "Learn more" link. A decorative graphic on the right side depicts a hand holding a shield, a padlock, and a magnifying glass over a document.

# Architecture

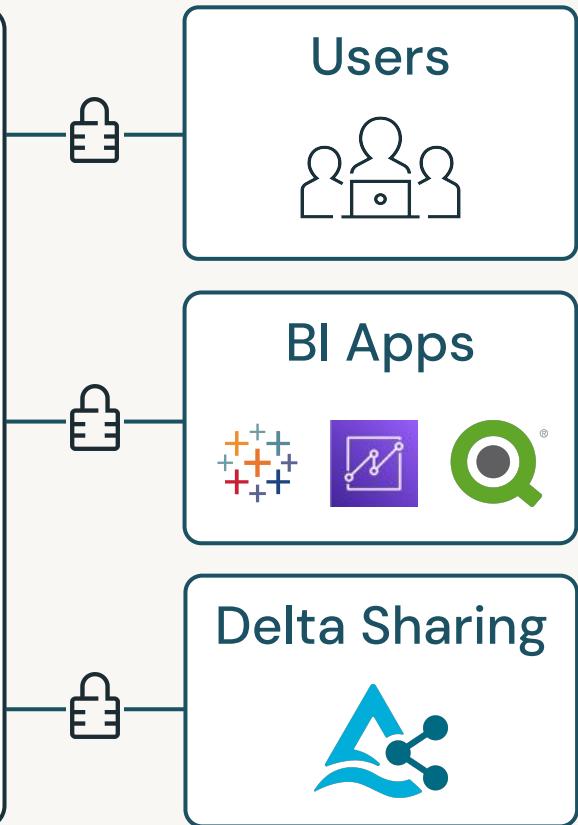
## Storage



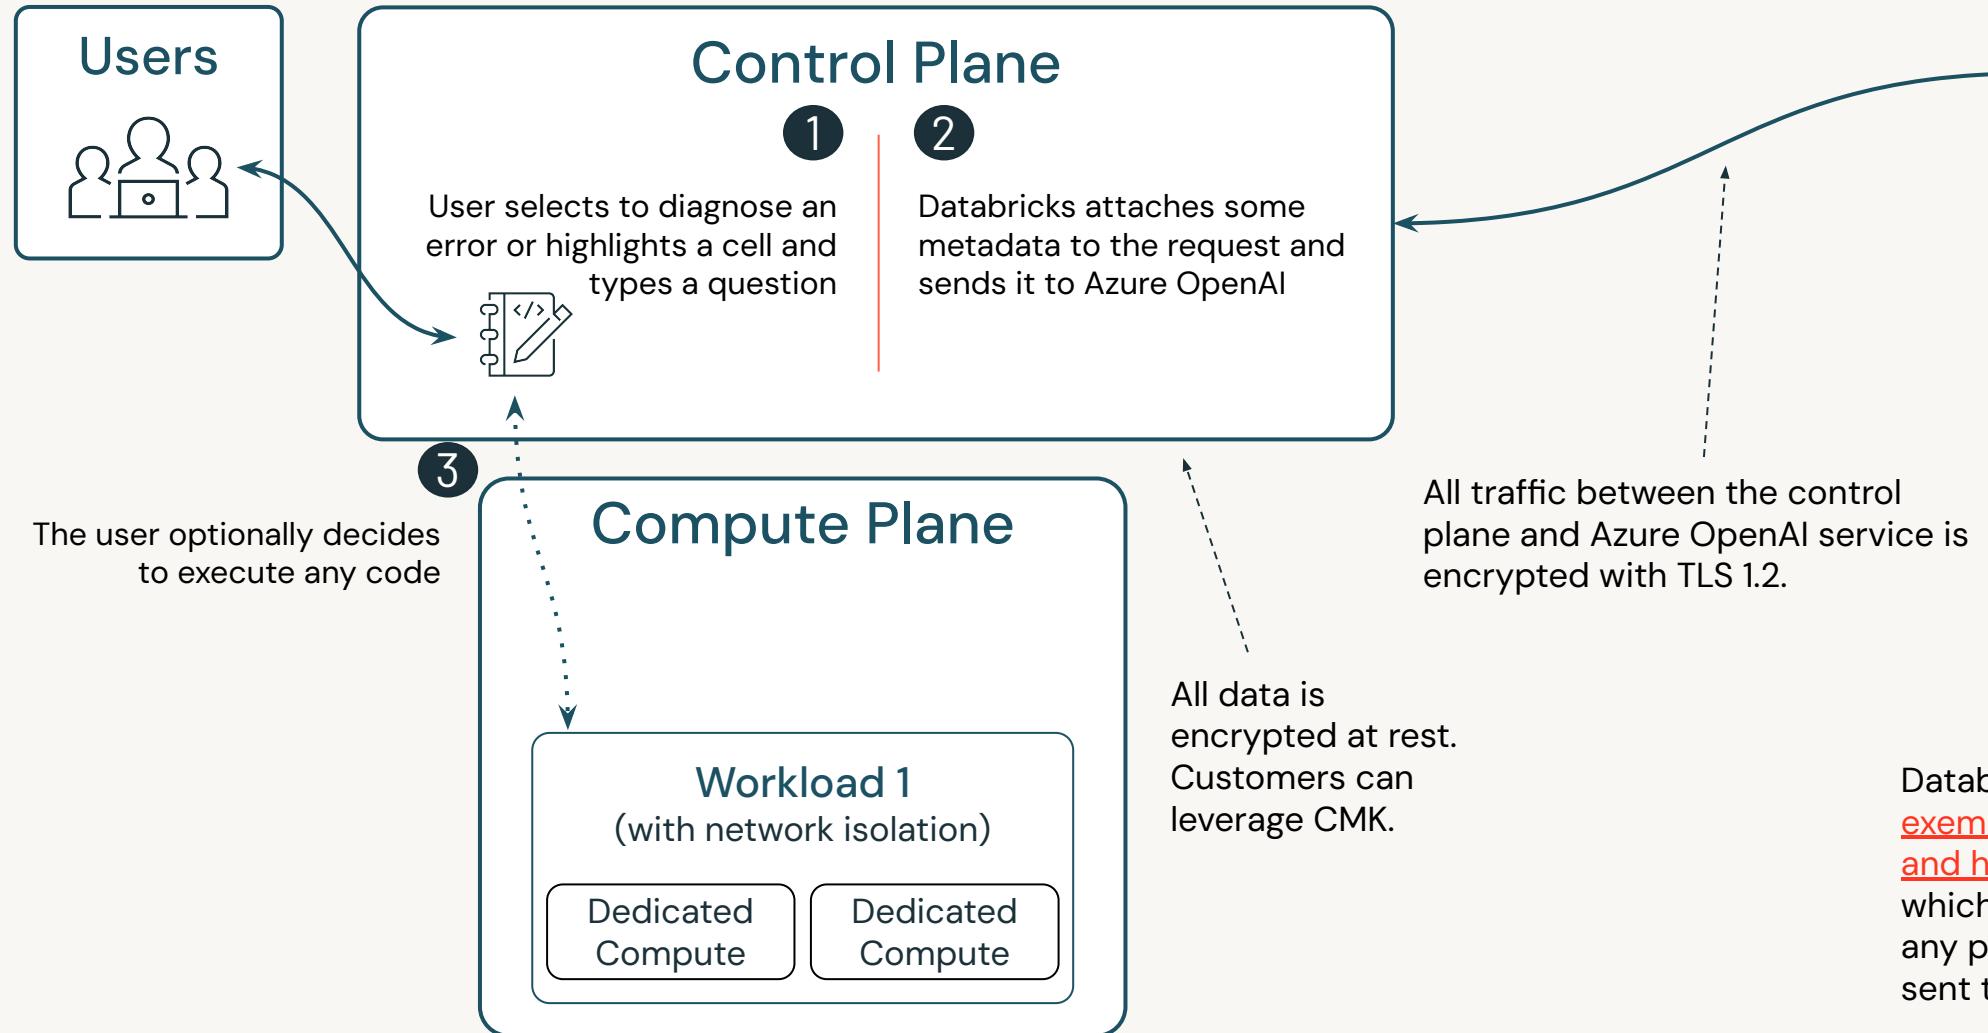
## Your Databricks Workspace



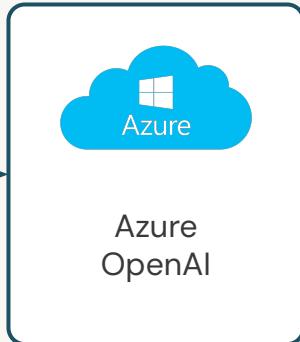
## Users



# Using the Databricks Assistant



Databricks uses different geographic-based Azure OpenAI services based on the geographic location of the workspace



Databricks has opted into [the exemption from abuse monitoring and human review](#) program, under which Microsoft does not store any prompts and completions sent to the Azure OpenAI service.

# Data FAQs

- **What data is being sent?**
  - Code or queries in the current notebook cell or SQL tab
  - Table and column names and descriptions
  - Previous prompt questions
  - Favourite tables
  - The “diagnose error” feature also shares the stack trace from the error output
  - We do **not** send row-level data
- **Does Azure OpenAI collect my data?**
  - No. Databricks has opted into the [exemption from abuse monitoring and human review](#) program, under which Microsoft does not store any prompts and completions sent to the Azure OpenAI service.
- **Are there any data residency considerations I should be aware of?**
  - We are currently using an Azure OpenAI service deployed in West Europe for all workspaces deployed within a European Geo and one in East US for everyone else. We will continue to evaluate support for other Azure OpenAI regions for future versions to meet latency and data residency requirements.

Please see the [DatabricksIQ trust and safety](#) page for more information

# Model FAQs

- **What models are you using?**
  - The Databricks Assistant is currently using Azure OpenAI GPT-3.5-turbo and GPT-4-turbo depending on the region you're in. We're continually evaluating new models and services and may include these in future iterations of the Assistant.
- **Is my data being used to train models?**
  - No, neither Databricks nor our model partner (Azure OpenAI) trains models using customer data.
- **Is any content filtering applied to filter for harmful responses?**
  - Databricks also uses Azure OpenAI content filtering to protect users from harmful content. In addition, Databricks has performed an extensive evaluation with thousands of simulated user interactions to ensure that the protections put in place to protect against harmful content, jailbreaks, insecure code generation, and use of third-party copyright content are effective.

Please see the [DatabricksIQ trust and safety](#) page for more information

# Encryption

All of the following encrypted by default:

- All traffic between the user, the control plane, the compute plane & cloud APIs
- ***All traffic between the control plane and the Azure OpenAI service***
- All storage in the control plane and all attached disks

Optionally enable customer-managed keys (CMK):

- Code, queries, results & secrets
- ***Databricks assistant notebook and DB SQL chat history***
- DBFS root storage
- Data you store in your cloud provider

# Data Privacy

## **What model are you using?**

Azure OpenAI. Are analyzing other internal and external models

## **Is my data used for model training?**

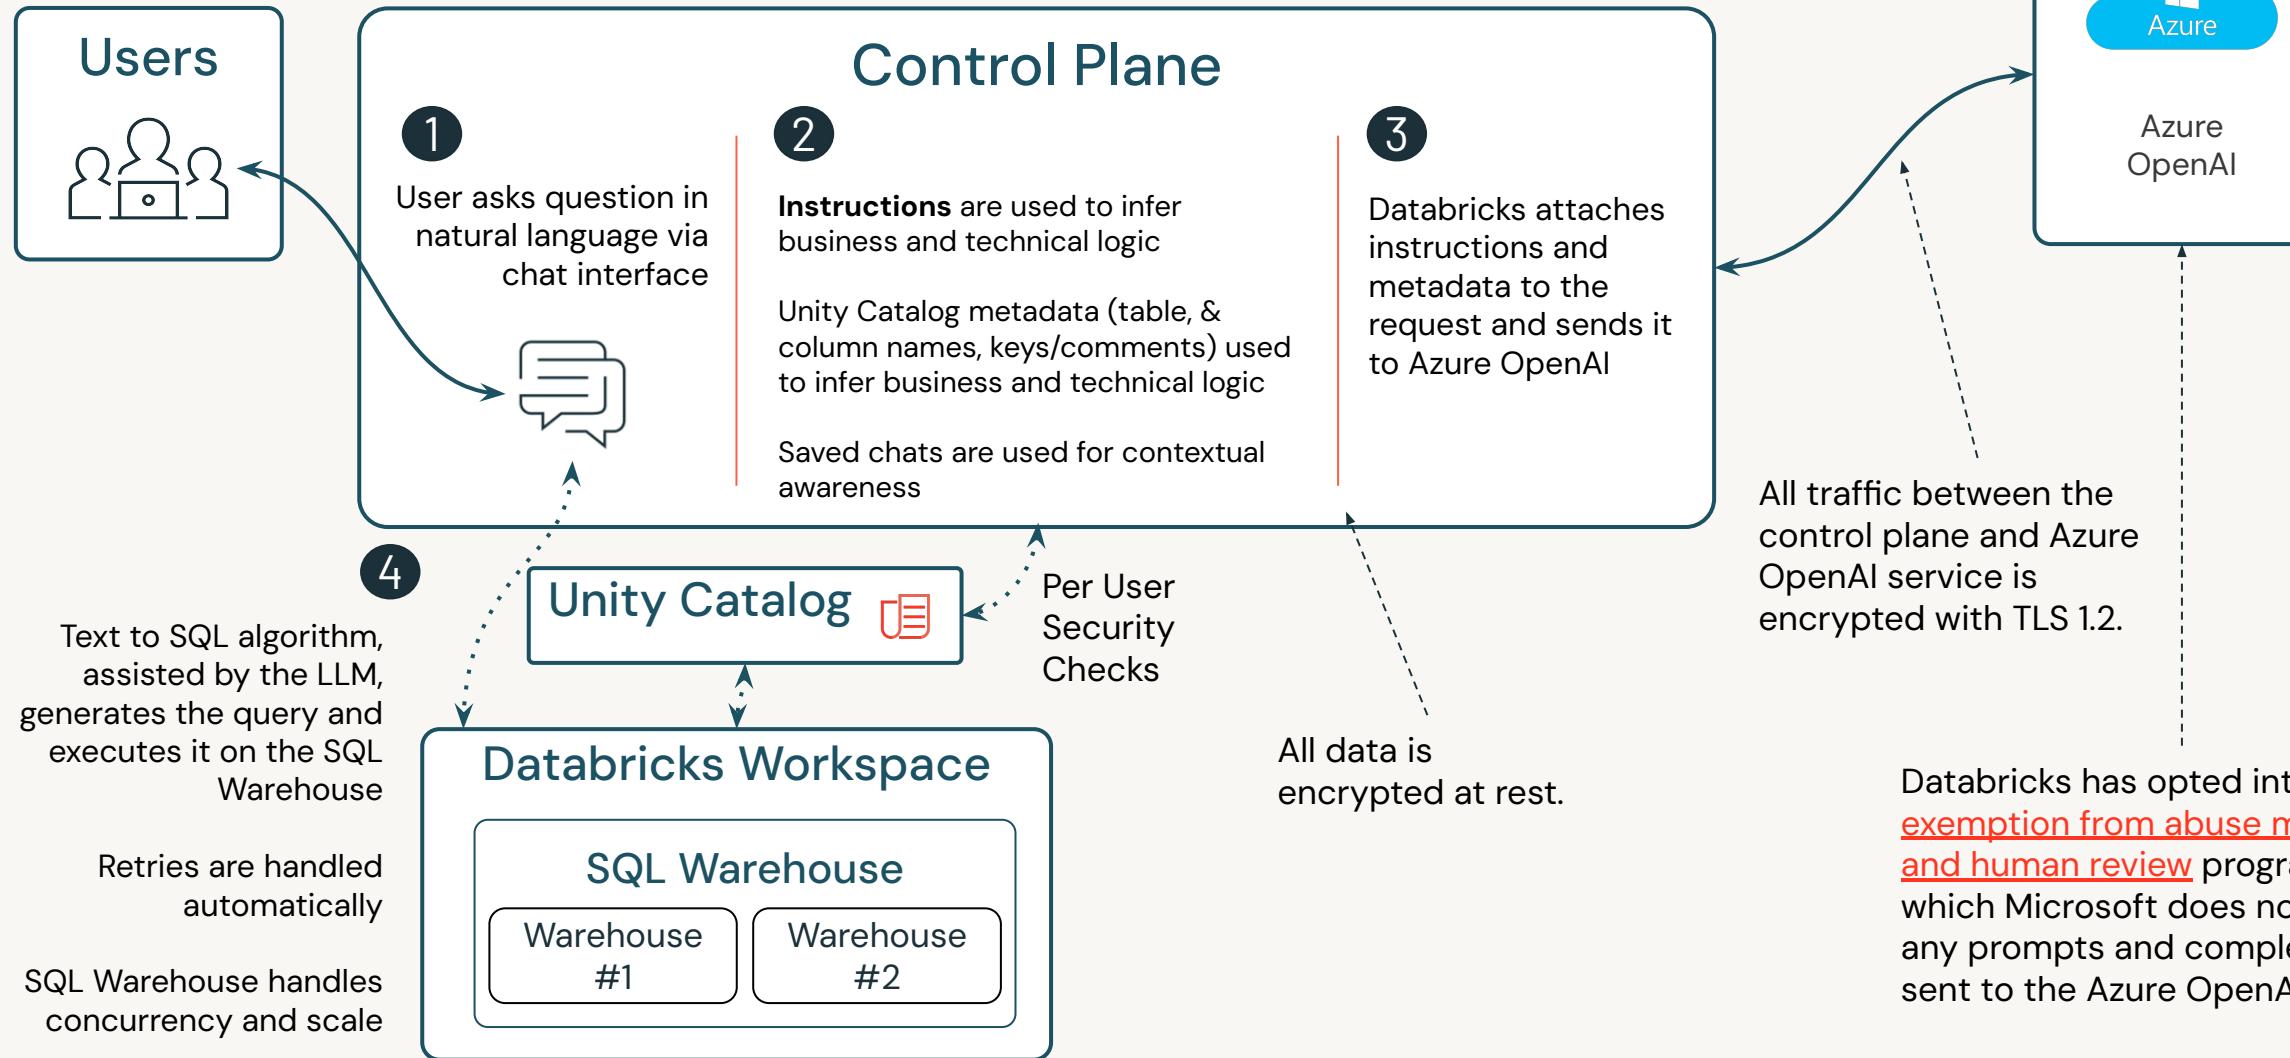
No. Neither Databricks nor OpenAI train models using customer data.

## **Do you send underlying customer data to the model?**

No. Values of specific rows are not sent.

Databricks uses different geographic-based Azure OpenAI services based on the geographic location of the workspace

# Data Rooms Architecture



Databricks has opted into [the exemption from abuse monitoring and human review](#) program, under which Microsoft does not store any prompts and completions sent to the Azure OpenAI service.



# Few things to note

- Currently in public preview and is free to use
- Final pricing model will be established and communicated prior to GA
- Admins can turn on/off this feature at account and workspace level

# References

## For further information

- [Demo Video](#)
- [Product Tour \(no login needed\)](#)
- [Add AI-generated Comments to Table](#)
- [Databricks Assistant FAQ](#)
- [Lakeview Dashboards Public Preview](#)
- [Lakeview Dashboards Overview](#)
- [5 tips to get the most out of your Databricks Assistant](#)
- [¿Qué es Databricks Assistant?](#)

**TRY**  
Databricks  
Assistant  
**TODAY**

Build your first Data+AI  
project in minutes

Take a Tour



PRODUCT TOUR 

Discover LakehouseIQ: The AI-Powered Engine That Uniquely Understands Your Business

[Launch Product Tour](#)

Watch a Demo



[Launch Demo Video](#)

# Demo

The code first creates the `average_car_prices` DataFrames using `pandas.read_csv()`. It then groups the DataFrames by the model column, calculates the average price of each model and stores the result in the `average_prices` DataFrame. Finally, the code selects the top 10 models and displays them using the `display` method. The `display` function is a specific method provided by Jupyter notebooks that allows you to display the output from the `print` function in a readable format.

<https://www.youtube.com/watch?v=q9REIGKkAVo&t=1s>

# Thank you



databricks